

# How to Debug Slow Spark Jobs?

Slow Spark jobs can burn through compute resources and frustrate engineers. Debugging performance issues requires a structured approach using **profiling** (Spark UI) and **logs** to pinpoint bottlenecks.

## 1. Understanding the Symptoms

Before diving into logs or tuning parameters, start by identifying performance symptoms:

- a. Long job execution time** (Jobs taking hours instead of minutes)
- b. Straggler tasks** (Some tasks significantly slower than others)
- c. High shuffle read/write** (Spilling data to disk, causing I/O overhead)
- d. Memory consumption issues** (Out-of-memory errors or excessive garbage collection)
- e. Driver bottlenecks** (High CPU usage on the driver node).

## Case Study: EMR to Databricks Migration — Job Running 4x Slower

We were helping a retail client **migrate their ETL workloads from AWS EMR to Databricks**. The job, which processed **1 billion transactions daily**, ran in **30 minutes on EMR** but **took 2+ hours on Databricks**.

The immediate question: **Why was the job suddenly 4x slower?**

## 2. Identifying the Bottlenecks

To debug Spark performance, break it down into:

### A. Task Execution Time

If **some tasks take much longer**, it could indicate **data skew**.

If **all tasks are slow**, look at shuffle performance or inefficient transformations.

### B. Shuffle and Data Skew

**Large shuffle writes** can lead to **disk I/O bottlenecks**.

Uneven shuffle distribution can create **straggler tasks**.

### C. Memory and Garbage Collection

**Frequent task failures** → Check for **executor OOM errors**.

**Excessive GC time (>10% of job duration)** → Memory-intensive operations like large joins.

### D. Driver Bottlenecks

If the driver CPU is high, check for **collect()** or **large dataset broadcasting**.

## 3. Using Spark UI for Profiling in Databricks

**Databricks provides deep insights into job execution through the Spark UI.**  
Here's how we used it for debugging:

## A. Checking Stage Execution Timeline

### Databricks UI → Jobs → Completed Jobs → Stages

1. We found **one stage taking 70% of the total job time**.
2. The **task distribution was uneven**, with some tasks finishing in 2 minutes while others took 15 minutes.

**Root Cause: Severe data skew** in a key column used for joining fact and dimension tables.

## B. Analyzing DAG Visualization

The **DAG (Directed Acyclic Graph) Visualization** in Databricks UI revealed:

**Multiple wide transformations** (shuffle-heavy operations). **A large join operation with a skewed key** (causing straggler tasks).

**Fix:** We applied **salting to balance partitions** and used **broadcast joins** to reduce shuffle.

```
Copyfrom pyspark.sql.functions import monotonically_increasing_id
```

```
df = df.withColumn("salt", monotonically_increasing_id() % 10)
```

```
df_joined = df1.join(df2, ["common_column", "salt"])
```

## C. Monitoring Executors and Storage

We checked the **"Executors" tab in Spark UI** and found:

**Executors were frequently restarting** due to OOM errors. **Shuffle spill exceeded memory limits**, causing disk writes.

**Fix:** We optimized shuffle memory allocation:

```
Copyspark.sql.shuffle.partitions=200 # Reduce partitions from 2000 to 200
```

```
spark.memory.fraction=0.6 # Allocate more memory for execution
```

## 4. Digging into Logs for Root Cause Analysis in Databricks

Databricks provides **detailed logs** in the **Driver Logs (stderr/stdout)** and execution history.

### A. Checking Executor Logs for OOM Errors

We found **OutOfMemoryErrors** in the executor logs:

```
Copygrep -i "OutOfMemoryError" driver-logs.log
```

**Fix:** Increased executor memory:

```
Copyspark.executor.memory=6g
```

```
spark.executor.memoryOverhead=2g
```

### B. Debugging Shuffle Performance Issues

The shuffle logs indicated **excessive spill to disk**:

```
Copygrep -i "shuffle" stderr.log
```

**Fix:** Reduced shuffle partition size and used **broadcast joins**.

```
Copyfrom pyspark.sql.functions import broadcast
```

```
df = df1.join(broadcast(df2), "common_column")
```

### C. Identifying Long-Running Tasks

We ran a log analysis script:

```
Copygrep -i "task" executor-logs.log | sort -nk3 | tail -10
```

**Fix:** Repartitioned based on skewed keys.

```
Copydf = df.repartition(100, "key_column")
```

## 5. Final Optimizations and Performance Gains

After applying **targeted fixes**, the job execution time **dropped from 2+ hours to 28 minutes**, slightly better than EMR.

**Key Improvements:**

Balanced partitions to avoid skew  
Reduced shuffle operations using broadcast joins  
Increased executor memory to handle large datasets  
Optimized memory management to reduce GC time

## Conclusion

Debugging slow Spark jobs in **Databricks** requires:

1. **Using Spark UI** to identify slow stages and tasks.
2. **Checking logs** for OOM errors, shuffle performance, and executor restarts.
3. **Applying optimizations** like repartitioning, broadcasting, and memory tuning.

By following these steps, **80% of performance issues** can be resolved with **minimal tuning**.